

Homework #4: Boosting, Dimension reduction, Clustering

We will analyze data from the Investigation of Serial Studies to Predict Your Therapeutic Response with Imaging and moLecular Analysis (I-SPY TRIAL) breast cancer trial. This dataset contains patients with measurements derived from MRI scans. The outcome of interest is to predict the final diameter of the tumor as measured through the MRI, which is stored in the column `MRI_LD_Tfinal`. We have longitudinal measurements from each patient, collected at timepoints T0, T1, T2, and Tfinal. T0 is pre-treatment. T1 is early treatment. T2 is inter-regimen. Tfinal is the final measurement.

1. Load the iSPY1 dataset by running the following. Notice that there is a `fakeSite` column with values ranging from 1-8. We've added this column to simulate combining data from 8 different sites.

```
ispy_dat <- read.csv("ispy1doctored_site.csv")
ispy_dat$HR_HER2status <- as.factor(ispy_dat$HR_HER2status)
```

Boosting

We will create and evaluate a gradient boosted model that predicts `MRI_LD_Tfinal` using all the predictors measured before Tfinal.

2. (1 point) Hold out sites 7 and 8 for testing. Store the training data in `datTrain` and the test data in `datTest`.

```
datTest <- ispy_dat[ispy_dat["fakeSite"]==7|ispy_dat["fakeSite"]==8,]
datTrain <- ispy_dat[ispy_dat["fakeSite"]!=7&ispy_dat["fakeSite"]!=8,]
```

3. (2 points) In this homework, we will manually implement site-wise 3-fold cross-validation (i.e. two sites per fold) rather than using the `caret` package. Create vector with True and False values to split the rows in `datTrain` into 3 folds, where sites 1 and 2 are in the first fold, sites 3 and 4 are in the second fold, and sites 5 and 6 are in the third fold. Create a list with these three True/False vectors. It may be helpful to reference this list later in this homework.

```
cv1 <- datTrain$fakeSite %in% c(1,2)
cv2 <- datTrain$fakeSite %in% c(3,4)
cv3 <- datTrain$fakeSite %in% c(5,6)

cv_list <- list(as.numeric(cv1),as.numeric(cv2),as.numeric(cv3))
```

4. Load the `gbm` package.

```
library(gbm)
```

```
## Warning: package 'gbm' was built under R version 4.3.2
```

```
## Loaded gbm 2.1.9
```

This version of gbm is no longer under development. Consider transitioning to gbm3, <https://github.com/gbm-developers/gbm3>

5. (3 points) We will select the hyperparameters for a gradient boosted model using site-wise three-fold CV. Create a function named `fit_fold` that takes as input the fold number `fold_idx`, number of trees, and interaction depth. The function will fit a gradient boosted model using `gbm` using the aforementioned predictors. Also exclude `fakeSite`. Train on all the folds except for the `fold_idx`-th one. The function should output the mean squared error of the fitted model on the held out fold. Use the folds you made in question 3. Fix the shrinkage hyperparameter in `gbm` as 0.01.

```
fit_fold <- function(fold_idx, n_trees, interaction_depth){
  X <- datTrain[!(names(datTrain) %in% c('fakeSite', 'MRI_LD_Tfinal'))]
  y <- datTrain$MRI_LD_Tfinal

  train_row <- !cv_list[[fold_idx]]
  test_row <- cv_list[[fold_idx]]

  train_data <- X[train_row,]
  test_data <- X[test_row,]

  train_outcome <- y[train_row]
  test_outcome <- y[test_row]

  gbm_md1 <- gbm(
    formula = train_outcome ~ .,
    data = train_data,
    distribution = "gaussian",
    n.trees = n_trees,
    interaction.depth = interaction_depth,
    shrinkage = 0.01,
    n.minobsinnode = 10,
    verbose = FALSE
  )

  pred <- predict(gbm_md1, newdata= test_data, n.trees = n_trees)

  mse <- mean((test_outcome - pred)^2)
  return(mse)
}
```

6. (4 points) Using the function you made in question 5, tune the number of trees and interaction depth using site-wise three-fold CV. Search over the values `n.trees=100, 200, 400, 800, 1600` and `interaction.depth=1, 2`. Which hyperparameter values minimize the cross-validated mean squared error?

```

# Define the sets of potential hyperparameters
set.seed(7)
n_trees_values <- c(100, 200, 400, 800, 1600)
interaction_depth_values <- c(1, 2)

# Initialize a variable to store the results
best_mse <- Inf
results <- data.frame(n_trees = integer(), interaction_depth = integer(), avg_mse = numeric())
best_params <- list(n_trees = NA, interaction_depth = NA)

# Evaluate all combinations of hyperparameters
for (n_trees in n_trees_values) {
  for (interaction_depth in interaction_depth_values) {
    # Initialize a vector to store MSEs
    mse_values <- numeric(length(cv_list))

    # Perform 3-fold CV
    for (fold_idx in seq_along(cv_list)) {
      mse_values[fold_idx] <- fit_fold(fold_idx, n_trees, interaction_depth)
    }

    # Calculate the average MSE
    avg_mse <- mean(mse_values)
    results <- rbind(results, data.frame(n_trees = n_trees, interaction_depth = interaction_depth, avg_mse = avg_mse))

    # Update best hyperparameters if the current combination has a lower average MSE
    if (avg_mse < best_mse) {
      best_mse <- avg_mse
      best_params <- list(n_trees = n_trees, interaction_depth = interaction_depth)
    }
  }
}

# Print the best hyperparameters and the Lowest MSE
print(paste("Best n_trees:", best_params$n_trees))

```

```
## [1] "Best n_trees: 1600"
```

```
print(paste("Best interaction_depth:", best_params$interaction_depth))
```

```
## [1] "Best interaction_depth: 2"
```

```
print(paste("Lowest cross-validated MSE:", best_mse))
```

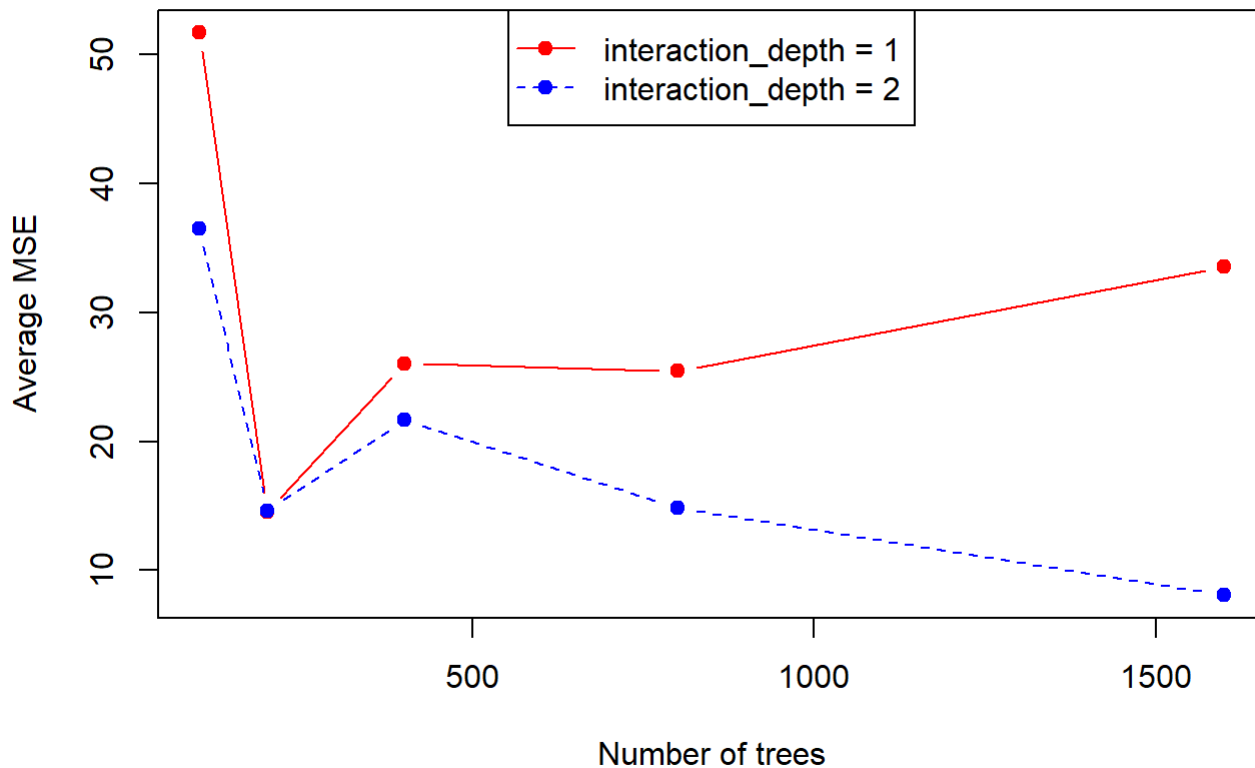
```
## [1] "Lowest cross-validated MSE: 8.07412546170622"
```

7. (2 points) Plot the cross-validated error with respect to `n.trees` for the interaction depth that attained the lowest CV error.

```
plot(results$n_trees[results$interaction_depth==1], results$avg_mse[results$interaction_depth==1],
     type = "b", col="red", xlab="Number of trees", ylab = "Average MSE", xlim = c(min(results$n_trees), max(results$n_trees)), ylim = c(min(results$avg_mse), max(results$avg_mse)), pch=19, lty=1)

points(results$n_trees[results$interaction_depth == 2], results$avg_mse[results$interaction_depth == 2],
       type = "b", col = "blue", pch = 19, lty = 2)

legend("top", legend = c("interaction_depth = 1", "interaction_depth = 2"),
      col = c("red", "blue"), pch = 19, lty = 1:2)
```



8. (1 point) Refit the gradient boosted model on all the training data (`datTrain`) using the hyperparameters that minimized the CV error.

```
# Check the best parameters and set them
opt_n_trees <- 1600
opt_interaction_depth <- 2


# Set the target and features
X <- datTrain[, !(names(datTrain) %in% c('fakeSite', 'MRI_LD_Tfinal'))]
y <- datTrain$MRI_LD_Tfinal

# Fit the model on the all train data
best_gbm_model <- gbm(
  formula = y ~ .,
  data = X,
  distribution = "gaussian",
  n.trees = opt_n_trees,
  interaction.depth = opt_interaction_depth,
  shrinkage = 0.01,
  n.minobsinnode = 10,
  verbose = TRUE
)
```

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	806.4364	nan	0.0100	5.1388
##	2	796.6885	nan	0.0100	8.5387
##	3	786.8402	nan	0.0100	9.6934
##	4	777.5080	nan	0.0100	9.0974
##	5	770.1237	nan	0.0100	5.3455
##	6	760.8756	nan	0.0100	8.0408
##	7	753.1938	nan	0.0100	7.4499
##	8	743.4184	nan	0.0100	7.2395
##	9	735.5992	nan	0.0100	6.3347
##	10	729.5958	nan	0.0100	5.8279
##	20	653.9902	nan	0.0100	5.1653
##	40	542.4794	nan	0.0100	2.4892
##	60	453.6630	nan	0.0100	2.9926
##	80	391.8812	nan	0.0100	2.0529
##	100	347.6759	nan	0.0100	0.7893
##	120	311.9355	nan	0.0100	1.1754
##	140	290.0601	nan	0.0100	-0.1587
##	160	270.7253	nan	0.0100	1.0622
##	180	253.2096	nan	0.0100	0.4521
##	200	240.3764	nan	0.0100	-0.0919
##	220	228.5378	nan	0.0100	-0.9384
##	240	218.7614	nan	0.0100	0.0132
##	260	211.2403	nan	0.0100	0.1441
##	280	204.3398	nan	0.0100	-0.5538
##	300	199.8679	nan	0.0100	-0.4728
##	320	194.3170	nan	0.0100	-0.4069
##	340	188.7469	nan	0.0100	-0.3721
##	360	184.7131	nan	0.0100	-0.1238
##	380	180.7707	nan	0.0100	-0.2276
##	400	175.8266	nan	0.0100	-0.1436
##	420	172.1787	nan	0.0100	-0.2954
##	440	168.2319	nan	0.0100	0.1000
##	460	164.9999	nan	0.0100	-0.4498
##	480	161.7935	nan	0.0100	-0.1838
##	500	159.2395	nan	0.0100	-0.1074
##	520	155.9049	nan	0.0100	-0.1099
##	540	153.1775	nan	0.0100	-0.4268
##	560	150.2953	nan	0.0100	-0.2861
##	580	147.8680	nan	0.0100	-0.4435
##	600	144.6684	nan	0.0100	-0.1660
##	620	142.3697	nan	0.0100	-0.3934
##	640	140.2654	nan	0.0100	-0.3305
##	660	138.2703	nan	0.0100	-0.0986
##	680	136.3267	nan	0.0100	-0.2439
##	700	134.2245	nan	0.0100	-0.3809
##	720	131.6922	nan	0.0100	-0.3878
##	740	129.6761	nan	0.0100	-0.2734
##	760	127.8547	nan	0.0100	-0.2613
##	780	125.9649	nan	0.0100	-0.0802
##	800	123.7761	nan	0.0100	-0.2306
##	820	121.7287	nan	0.0100	-0.0667

##	840	119.9784	nan	0.0100	-0.2200
##	860	118.2015	nan	0.0100	-0.1626
##	880	116.5736	nan	0.0100	-0.2371
##	900	114.5591	nan	0.0100	-0.1656
##	920	112.9851	nan	0.0100	-0.1196
##	940	111.2883	nan	0.0100	-0.2304
##	960	109.5806	nan	0.0100	-0.1310
##	980	107.8985	nan	0.0100	-0.2377
##	1000	106.5801	nan	0.0100	-0.1231
##	1020	104.7583	nan	0.0100	-0.3479
##	1040	103.0241	nan	0.0100	-0.0623
##	1060	101.7254	nan	0.0100	-0.1908
##	1080	100.0340	nan	0.0100	-0.1160
##	1100	98.7143	nan	0.0100	-0.1534
##	1120	97.4235	nan	0.0100	-0.2574
##	1140	96.5062	nan	0.0100	-0.2059
##	1160	95.1669	nan	0.0100	-0.2085
##	1180	93.9834	nan	0.0100	-0.1881
##	1200	92.9486	nan	0.0100	-0.2596
##	1220	91.8643	nan	0.0100	-0.2631
##	1240	90.4865	nan	0.0100	-0.1568
##	1260	89.3648	nan	0.0100	-0.2702
##	1280	88.1566	nan	0.0100	-0.1274
##	1300	86.9108	nan	0.0100	-0.1341
##	1320	85.8337	nan	0.0100	-0.0220
##	1340	84.8670	nan	0.0100	-0.2727
##	1360	83.6808	nan	0.0100	-0.0818
##	1380	82.5478	nan	0.0100	-0.1121
##	1400	81.6271	nan	0.0100	-0.0705
##	1420	79.9900	nan	0.0100	-0.2341
##	1440	79.1057	nan	0.0100	-0.1325
##	1460	78.3485	nan	0.0100	-0.3535
##	1480	77.5392	nan	0.0100	-0.1955
##	1500	76.4349	nan	0.0100	-0.1368
##	1520	75.5284	nan	0.0100	-0.1067
##	1540	74.5213	nan	0.0100	-0.2209
##	1560	73.4533	nan	0.0100	-0.1799
##	1580	72.6055	nan	0.0100	-0.1235
##	1600	71.4229	nan	0.0100	-0.2966

```
summary(best_gbm_model)
```



```
##          var      rel.inf
## MRI_LD_T2      MRI_LD_T2 45.7025693
## ftvPeT2        ftvPeT2  8.1709222
## ftvPeT1        ftvPeT1  7.5837720
## ftvPeTfinal    ftvPeTfinal 6.5793962
## MRI_LD_T1      MRI_LD_T1  5.4166433
## MRI_LD_T0      MRI_LD_T0  4.5408821
## age            age      4.2300717
## ftvPeT0        ftvPeT0  3.6673401
## ftvPePctChgT0_T1 ftvPePctChgT0_T1 3.5291705
## HR_HER2status  HR_HER2status 3.2675751
## pixelVolT1     pixelVolT1  1.9706736
## pixelVolT0     pixelVolT0  1.4228526
## pixelVolTfinal pixelVolTfinal 1.4040006
## pixelVolT2     pixelVolT2  1.3099910
## pixelVolPctChgT0_T1 pixelVolPctChgT0_T1 0.9202991
## race           race      0.2838404
```

9. (1 point) Evaluate the MSE of the fitted model on the test data. How much of the variance have we explained using the GBM?


```
test_predictions <- predict(best_gbm_model, newdata = datTest[, !(names(datTest) %in% c('fakeSite', 'MRI_LD_Tfinal'))], n.trees = opt_n_trees)

# Observed data in the test dataset
test_actual <- datTest$MRI_LD_Tfinal

# Calculate MSE and R^2 score
mse_test <- mean((test_actual - test_predictions)^2)
print(paste("Test MSE:", mse_test))
```

```
## [1] "Test MSE: 383.480561993256"
```

```
ss_total <- sum((test_actual - mean(test_actual))^2)
ss_res <- sum((test_actual - test_predictions)^2)
r_squared <- 1 - (ss_res / ss_total)
print(paste("R^2 Score:", r_squared))
```

```
## [1] "R^2 Score: 0.534967769074254"
```

Kmeans

Let's perform K-means on the iSPY data.

10. (1 point) Create a new data frame named `ispy_subdat` that only contains the continuous variables measured before Tfinal.

```
drop_col <- c('age', 'race', 'HR_HER2status', 'fakeSite')
ispy_subdat <- ispy_dat[, !(names(ispy_dat) %in% drop_col)]
```

11. (1 point) Before we run K-means, center and scale all the variables so that they have mean 0 and variance 1 (hint: use the `scale` command).

```
ispy_subdat_scaled <- scale(ispy_subdat)
```

12. (3 points) Tune the number of clusters used in K-means. To do this, use the function `fviz_nbclust` from the `factoextra` library. The function `fviz_nbclust` determines and visualizes the optimal number of clusters using different methods (within cluster sums of squares, average silhouette and gap statistics). Plot the average silhouette with respect to the number of clusters by passing in the argument `method="silhouette"`. What is the optimal number of clusters according to the silhouette statistic?

```
library(factoextra)
```


```
## Warning: package 'factoextra' was built under R version 4.3.2
```

```
## Loading required package: ggplot2
```

```
## Warning: package 'ggplot2' was built under R version 4.3.2
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
```

```
fviz_nbclust(ispy_subdat, kmeans, method = "silhouette")
```




The optimal number of clusters is two.

13. (3 points) Refit k-means using the optimal number of clusters with 15 random initializations. Use the function `fviz_cluster()` to plot the clusters from K-means. Observations are represented by points in the plot, using principal components if $p > 2$. An ellipse is drawn around each cluster.

```
n_cluster <- 2
set.seed(42)
km <- kmeans(ispy_subdat, centers = n_cluster, nstart = 15)
fviz_cluster(km, data = ispy_subdat)
```


Cluster plot



13. (1 point) Plot the distribution of `MRI_LD_Tfinal` for each cluster created by K-means. How do they differ across the clusters?

```
ispy_subdat$cluster <- km$cluster

boxplot(MRI_LD_Tfinal ~ cluster, data = ispy_subdat, xlab = "Cluster", ylab = "MRI_LD_Tfinal")
```



The cluster 1 is the subgroup of smaller MRI_LD_Tfinal and the cluster 2 is the bigger MRI_LD_Tfinal group.